

Churn Prediction

March 24, 2020

1 Predicting churn for bank

Importing all the necessary libraries

```
[200]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import QuantileTransformer
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.feature_selection import SelectFromModel
```

```
[201]: data = pd.read_csv("../input/predicting-churn-for-bank-customers/
↳Churn_Modelling.csv")
```

For checking correlation between features

```
[202]: data.corr()
```

```
[202]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	\
RowNumber	1.000000	0.004202	0.005840	0.000783	-0.006495	
CustomerId	0.004202	1.000000	0.005308	0.009497	-0.014883	
CreditScore	0.005840	0.005308	1.000000	-0.003965	0.000842	
Age	0.000783	0.009497	-0.003965	1.000000	-0.009997	
Tenure	-0.006495	-0.014883	0.000842	-0.009997	1.000000	
Balance	-0.009067	-0.012419	0.006268	0.028308	-0.012254	
NumOfProducts	0.007246	0.016972	0.012238	-0.030680	0.013444	
HasCrCard	0.000599	-0.014025	-0.005458	-0.011721	0.022583	
IsActiveMember	0.012044	0.001665	0.025651	0.085472	-0.028362	
EstimatedSalary	-0.005988	0.015271	-0.001384	-0.007201	0.007784	
Exited	-0.016571	-0.006248	-0.027094	0.285323	-0.014001	

	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
RowNumber	-0.009067	0.007246	0.000599	0.012044	
CustomerId	-0.012419	0.016972	-0.014025	0.001665	
CreditScore	0.006268	0.012238	-0.005458	0.025651	
Age	0.028308	-0.030680	-0.011721	0.085472	
Tenure	-0.012254	0.013444	0.022583	-0.028362	
Balance	1.000000	-0.304180	-0.014858	-0.010084	
NumOfProducts	-0.304180	1.000000	0.003183	0.009612	
HasCrCard	-0.014858	0.003183	1.000000	-0.011866	
IsActiveMember	-0.010084	0.009612	-0.011866	1.000000	
EstimatedSalary	0.012797	0.014204	-0.009933	-0.011421	
Exited	0.118533	-0.047820	-0.007138	-0.156128	

	EstimatedSalary	Exited
RowNumber	-0.005988	-0.016571
CustomerId	0.015271	-0.006248
CreditScore	-0.001384	-0.027094
Age	-0.007201	0.285323
Tenure	0.007784	-0.014001
Balance	0.012797	0.118533
NumOfProducts	0.014204	-0.047820
HasCrCard	-0.009933	-0.007138
IsActiveMember	-0.011421	-0.156128
EstimatedSalary	1.000000	0.012097
Exited	0.012097	1.000000

```
[203]: data.shape
```

```
[203]: (10000, 14)
```

Let's observe data

```
[204]: data.head()
```

```
[204]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	

```
4          2  125510.82          1          1          1
```

```
EstimatedSalary  Exited
0          101348.88      1
1          112542.58      0
2          113931.57      1
3           93826.63      0
4           79084.10      0
```

```
[205]: data["Geography"].unique() #checking for unique values in Geography
```

```
[205]: array(['France', 'Spain', 'Germany'], dtype=object)
```

```
[206]: data.describe()
```

```
[206]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure \
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000

	Balance	NumOfProducts	HasCrCard	IsActiveMember \
count	10000.000000	10000.000000	10000.00000	10000.000000
mean	76485.889288	1.530200	0.70550	0.515100
std	62397.405202	0.581654	0.45584	0.499797
min	0.000000	1.000000	0.00000	0.000000
25%	0.000000	1.000000	0.00000	0.000000
50%	97198.540000	1.000000	1.00000	1.000000
75%	127644.240000	2.000000	1.00000	1.000000
max	250898.090000	4.000000	1.00000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

```
[207]: data.dtypes
```

```
[207]: RowNumber          int64
      CustomerId         int64
      Surname            object
      CreditScore        int64
      Geography          object
      Gender             object
      Age               int64
      Tenure             int64
      Balance            float64
      NumOfProducts      int64
      HasCrCard          int64
      IsActiveMember     int64
      EstimatedSalary    float64
      Exited             int64
      dtype: object
```

2 Data Visualization

```
[208]: plt.figure(figsize = (15,15))
      sns.catplot(x = 'Geography', kind = 'count', data = data, palette = 'pink')
      plt.title('Customers distribution across Countries')
      plt.show()
```


<Figure size 1080x1080 with 0 Axes>



Maximum customers from France

```
[209]: plt.figure(figsize = (15,15))
sns.catplot(x = 'Gender', kind = 'count', data = data, palette = 'pastel')
plt.title("Males vs Females")
plt.show()
```

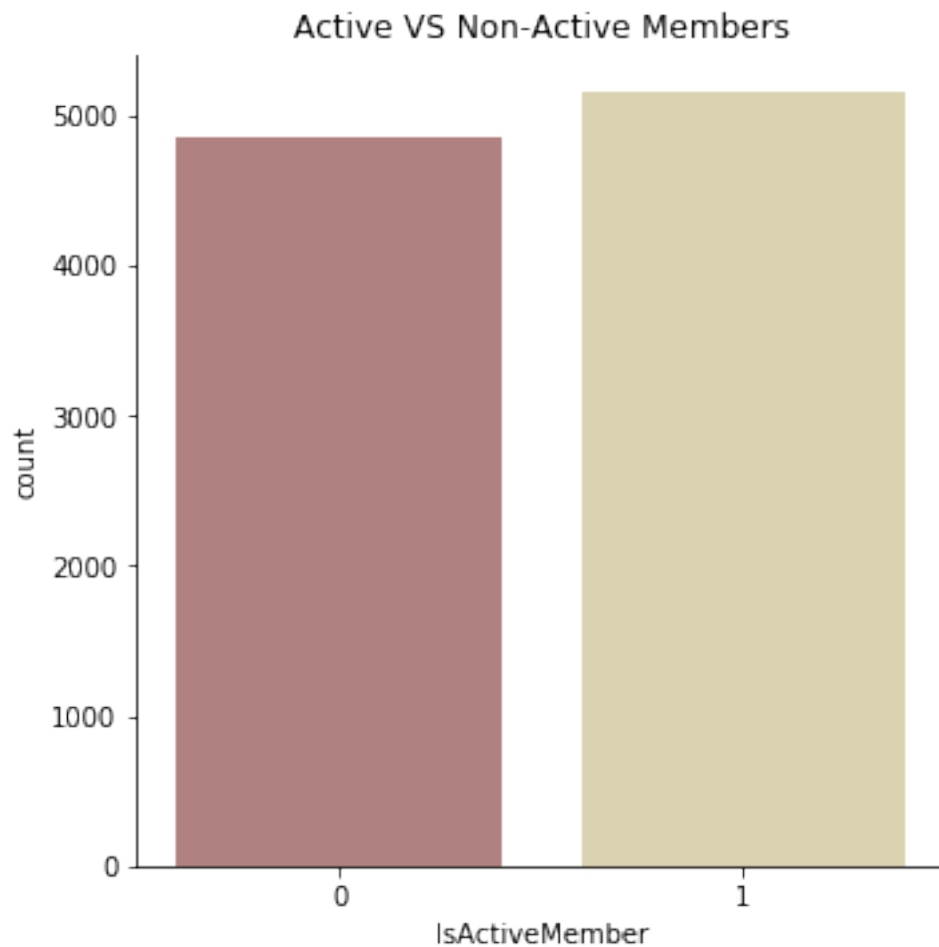
<Figure size 1080x1080 with 0 Axes>



We have more male customers

```
[210]: plt.figure(figsize = (15,15))
sns.catplot(x = 'IsActiveMember', kind = 'count', data = data, palette = 'pink')
plt.title("Active VS Non-Active Members")
plt.show()
```

<Figure size 1080x1080 with 0 Axes>



We have more active members

```
[211]: plt.figure(figsize = (15,15))
sns.catplot(x = 'HasCrCard', kind = 'count', palette = 'pastel', data = data)
plt.title("Credit Card VS No Credit Card")
plt.show()
```


<Figure size 1080x1080 with 0 Axes>



Most of the customers have credit card

```
[212]: plt.figure(figsize = (15,15))
sns.catplot(x = 'Exited', kind = 'count', hue = 'Gender', palette = 'pink',
↳data = data)
plt.title("Gender and Exited")
plt.show()
```


<Figure size 1080x1080 with 0 Axes>



Females are more likely to exit

```
[213]: plt.figure(figsize = (15,15))
sns.catplot(x = 'HasCrCard', kind = 'count', hue = 'Gender', palette = 'pastel', data = data)
plt.title("Gender and Credit Card")
plt.show()
```

<Figure size 1080x1080 with 0 Axes>

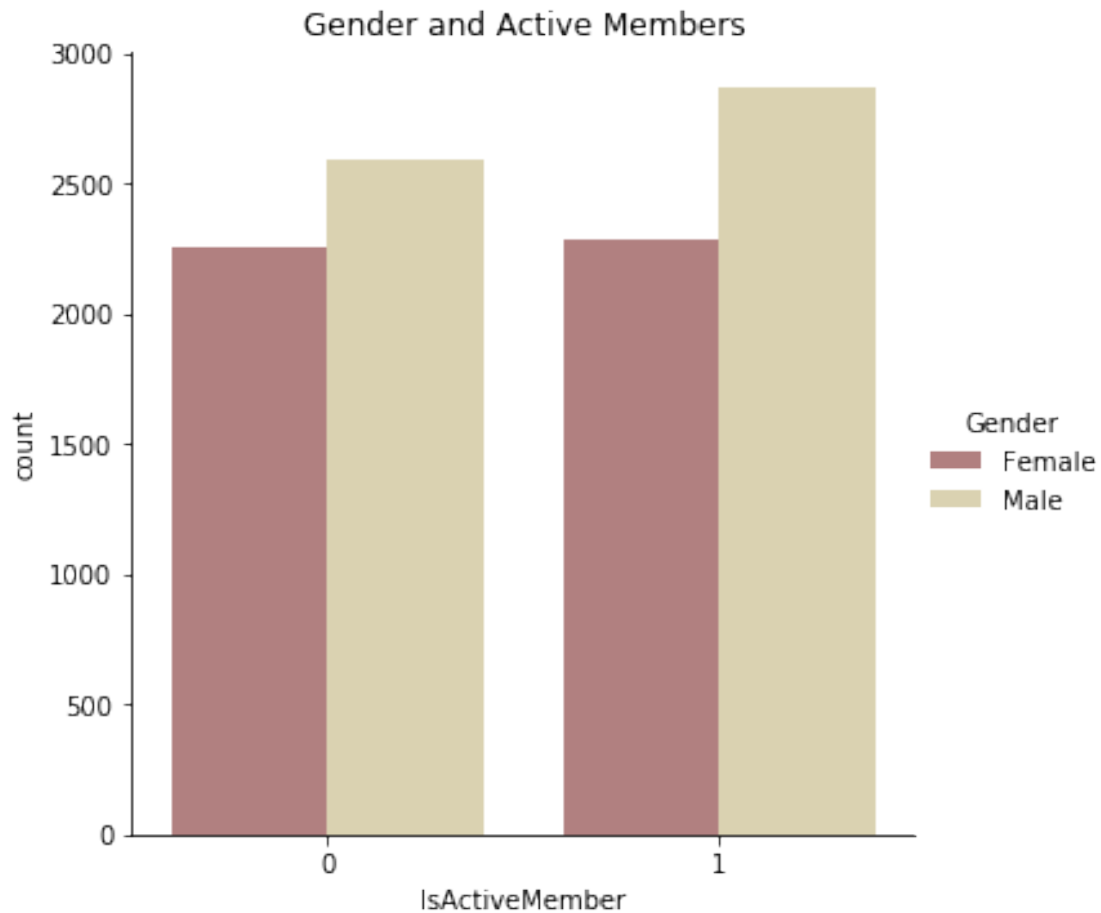


Males generally have credit card

But on other hand they are more likely not to have credit cards too

```
[214]: plt.figure(figsize = (15,15))
sns.catplot(x = 'IsActiveMember', kind = 'count', hue = 'Gender', palette = 'pink', data = data)
plt.title("Gender and Active Members")
plt.show()
```

<Figure size 1080x1080 with 0 Axes>




Males are more likely to be active members

But on other hand males are also likely to be non active members

```
[215]: plt.figure(figsize = (15,15))
sns.catplot(x = "NumOfProducts", kind = 'count', palette = 'pastel', data = data )
plt.title('Number of Products')
plt.show()
```

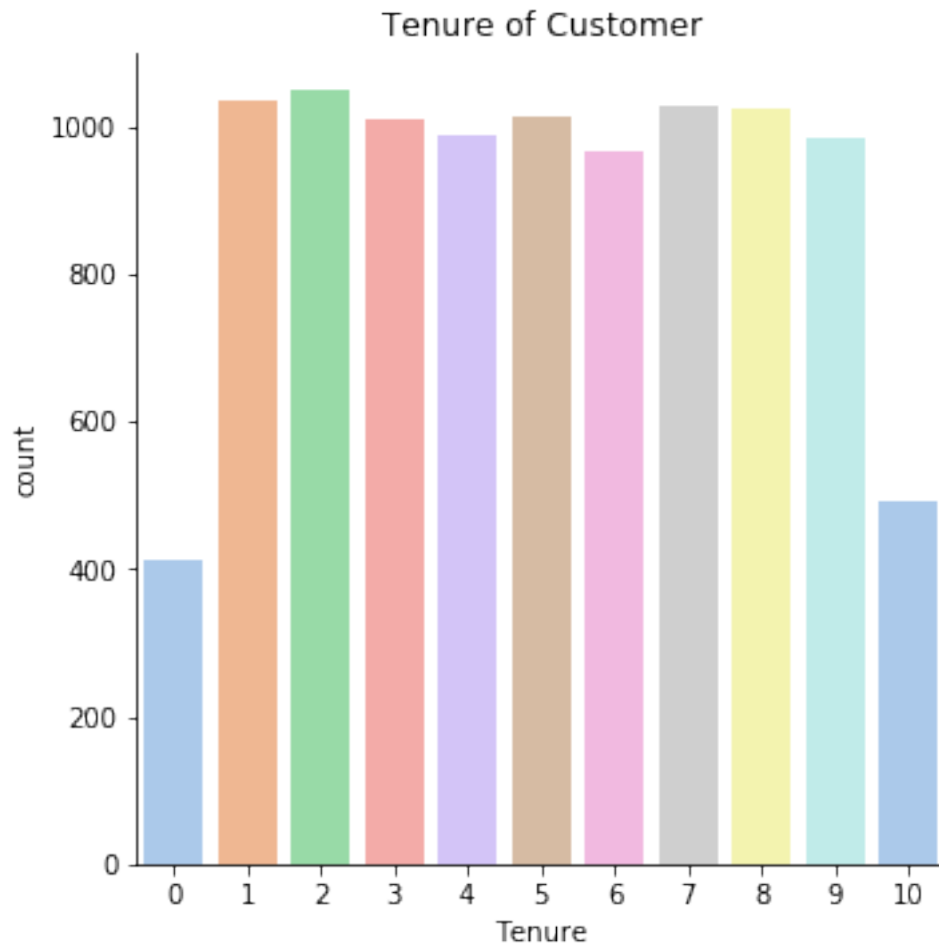
<Figure size 1080x1080 with 0 Axes>



Most of the customers have 1 or 2 products from bank

```
[216]: plt.figure(figsize = (15,15))
sns.catplot(x = 'Tenure', kind = 'count', palette = 'pastel', data = data)
plt.title("Tenure of Customer")
plt.show()
```


<Figure size 1080x1080 with 0 Axes>



Most customers have tenure of in between 1-8 years in bank

```
[217]: plt.figure(figsize = (15,15))
sns.catplot(x = 'Exited', kind = 'count', hue = 'IsActiveMember', palette = 'pink', data = data)
plt.title("Exited and Active Members")
plt.show()
```

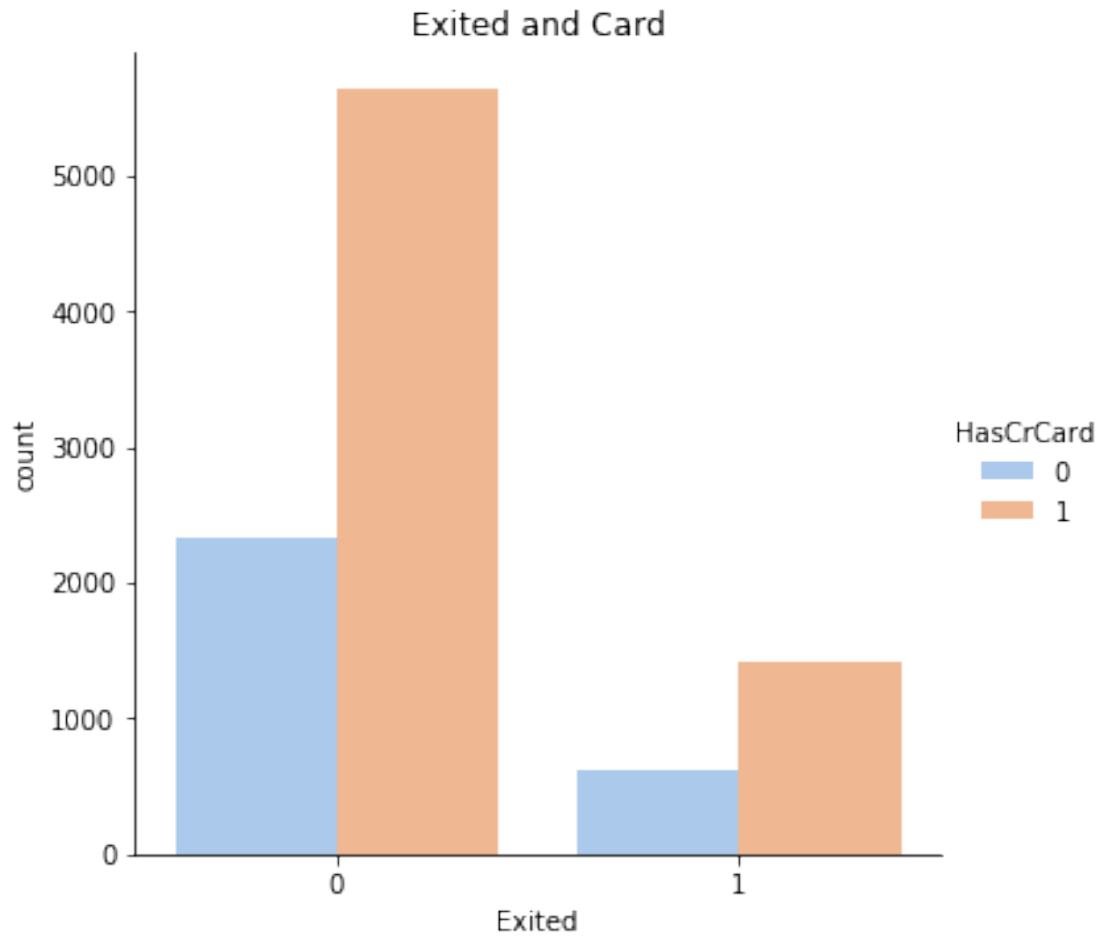
<Figure size 1080x1080 with 0 Axes>



Non active members are likely to exit more, quite understandable

```
[218]: plt.figure(figsize = (15,15))
sns.catplot(x = 'Exited', kind = 'count', hue = 'HasCrCard', palette = 'pastel', data = data)
plt.title("Exited and Card")
plt.show()
```

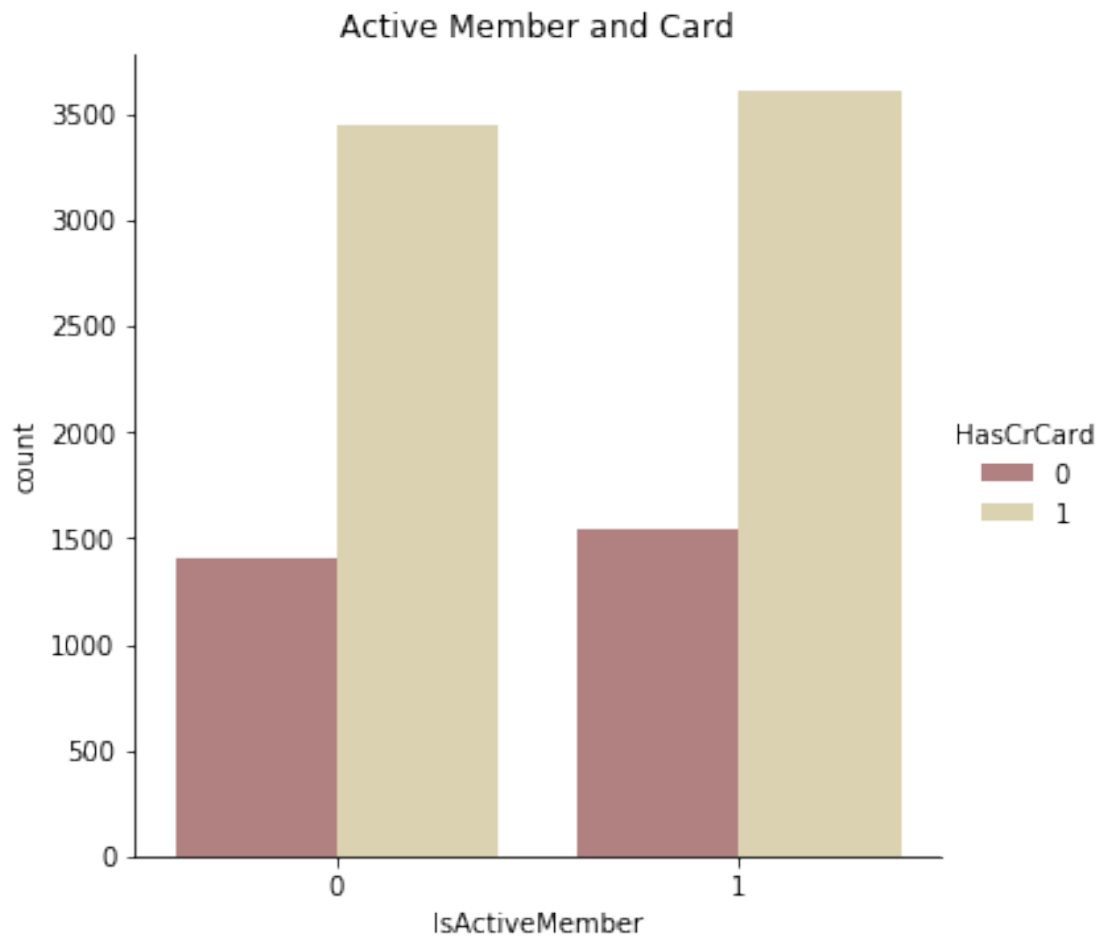
<Figure size 1080x1080 with 0 Axes>



Customers with credit card are likely to exit more

```
[219]: plt.figure(figsize = (15,15))
sns.catplot(x = 'IsActiveMember', kind = 'count', hue = 'HasCrCard', palette = 'pink', data = data)
plt.title('Active Member and Card')
plt.show()
```

<Figure size 1080x1080 with 0 Axes>



```
[220]: plt.figure(figsize = (15,15))
sns.scatterplot(x = 'Balance', y = 'EstimatedSalary', hue = 'Exited',palette = 'pastel', data = data)
plt.title("Balance vs Estimated Salary")
plt.show()
```




```
[221]: plt.figure(figsize = (15,15))
sns.scatterplot(x = 'Balance', y = 'CreditScore', hue = 'Exited',palette = _
↳ 'pink', data = data)
plt.title("Balance vs Credit Score")
plt.show()
```



```
[222]: plt.figure(figsize = (15,15))
sns.scatterplot(x = 'Balance', y = 'EstimatedSalary', hue = 'Gender',palette =_
↳'pastel', data = data)
plt.title("Estimated Salary vs Credit Score")
plt.show()
```



```
[223]: plt.figure(figsize = (15,15))
sns.scatterplot(x = 'Balance', y = 'EstimatedSalary', hue = 'Gender',
               ↪ 'IsActiveMember',palette = 'pastel', data = data)
plt.title("Estimated Salary vs Credit Score")
plt.show()
```



3 Data Preprocessing

```
[224]: data.drop(['RowNumber', 'CustomerId', 'Surname'], axis = 1, inplace = True)
```


```
[225]: data.isnull().sum() #checking for null values
```

```
[225]: CreditScore      0  
       Geography      0  
       Gender         0  
       Age            0  
       Tenure         0  
       Balance        0
```


```
NumOfProducts      0
HasCrCard           0
IsActiveMember      0
EstimatedSalary     0
Exited              0
dtype: int64
```

For checking skwness in the data

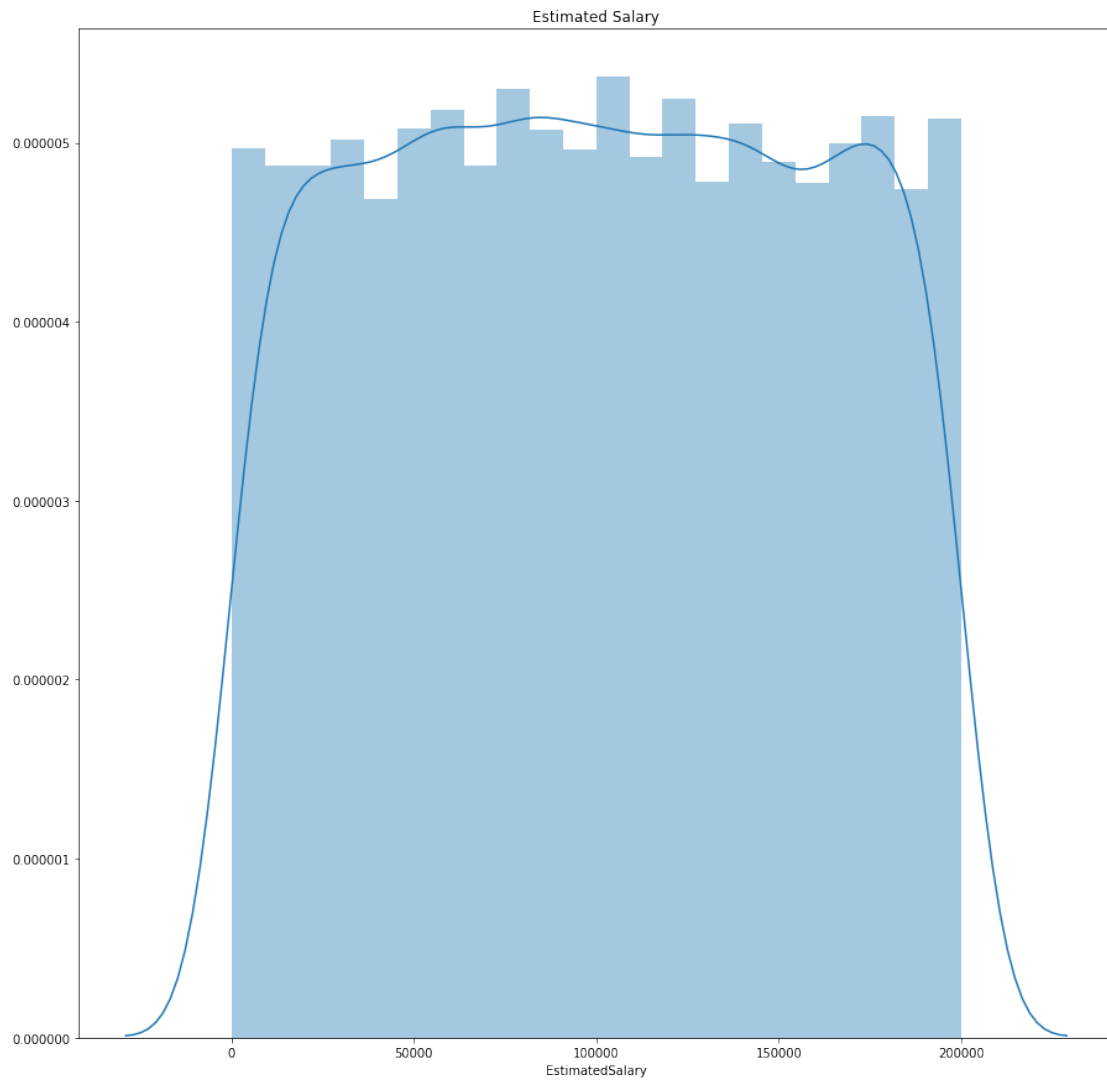
```
[226]: plt.figure(figsize = (15,15))
sns.distplot(data['Age'])
plt.title("Age")
plt.show()
```




```
[227]: plt.figure(figsize = (15,15))
sns.distplot(data["CreditScore"])
plt.title("Credit Score")
plt.show()
```



```
[228]: plt.figure(figsize = (15,15))
sns.distplot(data["EstimatedSalary"])
plt.title("Estimated Salary")
plt.show()
```





```
[229]: plt.figure(figsize = (15,15))
sns.distplot(data["Balance"])
plt.title("Balance")
plt.show()
```





For detecting outliers in data

```
[230]: column = ["Age", "Balance", "EstimatedSalary", "CreditScore"]
for i in column:
    plt.figure(figsize = (15,15))
    sns.boxplot(data[i])
    plt.title('Box Plot')
    plt.show()
```







```
[231]: data = data[(data["Age"] < 60)]
data = data[(data["CreditScore"] > 400)]
```

```
[232]: data.describe()
```

```
[232]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	\
count	9456.000000	9456.000000	9456.000000	9456.000000	9456.000000	
mean	650.898266	37.373308	5.018084	76521.194565	1.531514	
std	95.810805	8.316748	2.887855	62444.638692	0.579448	
min	401.000000	18.000000	0.000000	0.000000	1.000000	
25%	584.000000	31.000000	3.000000	0.000000	1.000000	
50%	652.000000	37.000000	5.000000	97302.205000	1.000000	

75%	717.000000	42.000000	8.000000	127644.240000	2.000000
max	850.000000	59.000000	10.000000	250898.090000	4.000000

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	9456.000000	9456.000000	9456.000000	9456.000000
mean	0.704949	0.500212	100176.595754	0.19797
std	0.456090	0.500026	57503.154035	0.39849
min	0.000000	0.000000	11.580000	0.00000
25%	0.000000	0.000000	51228.457500	0.00000
50%	1.000000	1.000000	100350.530000	0.00000
75%	1.000000	1.000000	149406.545000	0.00000
max	1.000000	1.000000	199992.480000	1.00000

Normalizing the data

```
[233]: data["Balance"] = QuantileTransformer().fit_transform(data["Balance"].values.
        ↪reshape(-1,1))
data["CreditScore"] = QuantileTransformer().fit_transform(data["CreditScore"].
        ↪values.reshape(-1,1))
data["EstimatedSalary"] = QuantileTransformer().
        ↪fit_transform(data["EstimatedSalary"].values.reshape(-1,1))
data["Age"] = QuantileTransformer().fit_transform(data["Age"].values.
        ↪reshape(-1,1))
```

```
[234]: data["Balance"] = StandardScaler().fit_transform(data["Balance"].values.
        ↪reshape(-1,1))
data["CreditScore"] = StandardScaler().fit_transform(data["CreditScore"].values.
        ↪reshape(-1,1))
data["EstimatedSalary"] = StandardScaler().fit_transform(data["CreditScore"].
        ↪values.reshape(-1,1))
```

```
[235]: data.describe()
```

```
[235]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	\
count	9.456000e+03	9456.000000	9456.000000	9.456000e+03	9456.000000	
mean	3.497156e-16	0.500092	5.018084	-1.539002e-16	1.531514	
std	1.000053e+00	0.288540	2.887855	1.000053e+00	0.579448	
min	-1.730136e+00	0.000000	0.000000	-1.211056e+00	1.000000	
25%	-8.629531e-01	0.229229	3.000000	-1.211056e+00	1.000000	
50%	7.675240e-04	0.512012	5.000000	1.829459e-01	1.000000	
75%	8.610264e-01	0.733734	8.000000	8.787345e-01	2.000000	
max	1.728209e+00	1.000000	10.000000	1.576003e+00	4.000000	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	9456.000000	9456.000000	9.456000e+03	9456.000000
mean	0.704949	0.500212	9.052263e-17	0.19797
std	0.456090	0.500026	1.000053e+00	0.39849

min	0.000000	0.000000	-1.730136e+00	0.000000
25%	0.000000	0.000000	-8.629531e-01	0.000000
50%	1.000000	1.000000	7.675240e-04	0.000000
75%	1.000000	1.000000	8.610264e-01	0.000000
max	1.000000	1.000000	1.728209e+00	1.000000

Label Encoding for categorical columns

```
[236]: data["Geography"] = LabelEncoder().fit_transform(data["Geography"])
data["Gender"] = LabelEncoder().fit_transform(data["Gender"])
```

```
[237]: data.head()
```

```
[237]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	-0.447536	0	0	0.733734	2	-1.211056	1	
1	-0.575623	2	0	0.697698	1	0.001502	1	
2	-1.505118	0	0	0.733734	8	1.411246	3	
3	0.632547	0	0	0.610110	1	-1.211056	2	
4	1.728209	2	0	0.766266	2	0.827109	1	

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	-0.447536	1
1	0	1	-0.575623	0
2	1	0	-1.505118	1
3	0	0	0.632547	0
4	1	1	1.728209	0

```
[238]: data.corr()
```

```
[238]:
```

	CreditScore	Geography	Gender	Age	Tenure	\
CreditScore	1.000000	0.008769	-0.004181	-0.010590	0.000090	
Geography	0.008769	1.000000	0.001382	0.035955	0.004042	
Gender	-0.004181	0.001382	1.000000	-0.032190	0.015198	
Age	-0.010590	0.035955	-0.032190	1.000000	-0.009384	
Tenure	0.000090	0.004042	0.015198	-0.009384	1.000000	
Balance	0.009128	0.064577	0.013669	0.041863	-0.013175	
NumOfProducts	0.009310	0.008042	-0.022634	-0.029477	0.014734	
HasCrCard	-0.003054	-0.012153	0.006354	-0.018187	0.020272	
IsActiveMember	0.022464	0.008056	0.020352	-0.011665	-0.026627	
EstimatedSalary	1.000000	0.008769	-0.004181	-0.010590	0.000090	
Exited	-0.017423	0.036878	-0.105182	0.345086	-0.012670	

	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
CreditScore	0.009128	0.009310	-0.003054	0.022464	
Geography	0.064577	0.008042	-0.012153	0.008056	
Gender	0.013669	-0.022634	0.006354	0.020352	
Age	0.041863	-0.029477	-0.018187	-0.011665	

Tenure	-0.013175	0.014734	0.020272	-0.026627
Balance	1.000000	-0.302849	-0.008085	-0.005267
NumOfProducts	-0.302849	1.000000	0.002371	0.008738
HasCrCard	-0.008085	0.002371	1.000000	-0.011784
IsActiveMember	-0.005267	0.008738	-0.011784	1.000000
EstimatedSalary	0.009128	0.009310	-0.003054	0.022464
Exited	0.113251	-0.048550	-0.009116	-0.137686

	EstimatedSalary	Exited
CreditScore	1.000000	-0.017423
Geography	0.008769	0.036878
Gender	-0.004181	-0.105182
Age	-0.010590	0.345086
Tenure	0.000090	-0.012670
Balance	0.009128	0.113251
NumOfProducts	0.009310	-0.048550
HasCrCard	-0.003054	-0.009116
IsActiveMember	0.022464	-0.137686
EstimatedSalary	1.000000	-0.017423
Exited	-0.017423	1.000000

4 Splitting Train and Test Data

```
[239]: y = data["Exited"]
```

```
[240]: y.head()
```

```
[240]: 0    1
       1    0
       2    1
       3    0
       4    0
       Name: Exited, dtype: int64
```

```
[241]: data.drop(["Exited"], axis = 1, inplace = True)
```

```
[242]: data.head()
```

```
[242]:   CreditScore  Geography  Gender    Age  Tenure  Balance  NumOfProducts  \
0   -0.447536         0      0  0.733734      2 -1.211056           1
1   -0.575623         2      0  0.697698      1  0.001502           1
2   -1.505118         0      0  0.733734      8  1.411246           3
3    0.632547         0      0  0.610110      1 -1.211056           2
4    1.728209         2      0  0.766266      2  0.827109           1
```

	HasCrCard	IsActiveMember	EstimatedSalary
0	1	1	-0.447536
1	0	1	-0.575623
2	1	0	-1.505118
3	0	0	0.632547
4	1	1	1.728209

```
[243]: train_x, test_x, train_y, test_y = train_test_split(data, y, test_size = 0.3,
↳ random_state = 50)
```

5 Model Fitting

6 Logistic Regression

```
[244]: logistic = LogisticRegression()
logistic.fit(train_x, train_y)
log_y = logistic.predict(test_x)
print(accuracy_score(log_y, test_y))
```

0.8290447655974621

Tuning the model

```
[245]: random_parameters = {'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000], 'penalty':
↳ ['l1', 'l2']}
print(random_parameters)
```

```
{'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000], 'penalty': ['l1', 'l2']}
```

```
[246]: random_para = RandomizedSearchCV(estimator = logistic, param_distributions =
↳ random_parameters, n_iter = 50, cv = 10, verbose=2, random_state= 50, n_jobs
↳ = -1)
random_para.fit(train_x, train_y)
```

Fitting 10 folds for each of 14 candidates, totalling 140 fits

```
/opt/conda/lib/python3.6/site-packages/sklearn/model_selection/_search.py:281:
UserWarning: The total space of parameters 14 is smaller than n_iter=50. Running
14 iterations. For exhaustive searches, use GridSearchCV.
```

```
% (grid_size, self.n_iter, grid_size), UserWarning)
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 38 tasks | elapsed: 2.1s
[Parallel(n_jobs=-1)]: Done 140 out of 140 | elapsed: 3.2s finished
```

```
[246]: RandomizedSearchCV(cv=10, error_score=nan,
estimator=LogisticRegression(C=1.0, class_weight=None,
```



```

        dual=False, fit_intercept=True,
        intercept_scaling=1,
        l1_ratio=None, max_iter=100,
        multi_class='auto', n_jobs=None,
        penalty='l2', random_state=None,
        solver='lbfgs', tol=0.0001,
        verbose=0, warm_start=False),
    iid='deprecated', n_iter=50, n_jobs=-1,
    param_distributions={'C': [0.001, 0.01, 0.1, 1, 10, 100,
                               1000],
                        'penalty': ['l1', 'l2']},
    pre_dispatch='2*n_jobs', random_state=50, refit=True,
    return_train_score=False, scoring=None, verbose=2)

```

```
[247]: random_para.best_params_
```

```
[247]: {'penalty': 'l2', 'C': 1}
```

```
[256]: logistic2 = LogisticRegression(penalty='l2', C=1)
logistic2.fit(train_x, train_y)
log_y = logistic2.predict(test_x)
print(accuracy_score(log_y, test_y))
```

0.8290447655974621

Feature Selection

```
[257]: feature = SelectFromModel(LogisticRegression())
feature.fit(train_x, train_y)
feature_support = feature.get_support()
feature_selected = train_x.loc[:, feature_support].columns.tolist()
print(str(len(feature_selected)), 'selected features')
```

2 selected features

```
[258]: print(feature_selected)
```

['Age', 'IsActiveMember']

```
[259]: train_x_feature = train_x[["Age", "IsActiveMember"]]
train_x_feature.head()
```

```
[259]:
```

	Age	IsActiveMember
2371	0.019520	1
6521	0.413914	0
9487	0.272773	1
2253	0.697698	1
2120	0.957958	1

```
[260]: test_x_feature = test_x[["Age", "IsActiveMember"]]
test_x_feature.head()
```

```
[260]:      Age  IsActiveMember
1624  0.056557            0
2169  0.766266            1
4285  0.655155            0
71    0.155155            0
4049  0.318318            1
```

```
[261]: logistic.fit(train_x_feature, train_y)
log_y_feature = logistic.predict(test_x_feature)
print(accuracy_score(log_y_feature, test_y))
```

0.8304547056750088

7 Random Forest Classifier

```
[262]: random = RandomForestClassifier()
random.fit(train_x,train_y)
random_y = random.predict(test_x)
print(accuracy_score(random_y,test_y))
```

0.8579485371871696

Tuning Parameters

```
[263]: n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
max_features = ['auto','sqrt']
max_depth = [int(x) for x in np.linspace(10,110,num=11)]
max_depth.append(None)
min_samples_split = [2,5,10]
min_samples_leaf = [1,2,4]
bootstrap = [True, False]
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap
              }
print(random_grid)
```

```
{'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000],
 'max_features': ['auto', 'sqrt'], 'max_depth': [10, 20, 30, 40, 50, 60, 70, 80,
 90, 100, 110, None], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2,
 4], 'bootstrap': [True, False]}
```

```
[264]: random_para = RandomizedSearchCV(estimator = random, param_distributions =_
    ↪random_grid, n_iter = 100, cv = 3, verbose=2, random_state=42, n_jobs = -1)
random_para.fit(train_x,train_y)
```

Fitting 3 folds for each of 100 candidates, totalling 300 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
/opt/conda/lib/python3.6/site-
packages/joblib/externals/loky/process_executor.py:706: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.
```

"timeout or by a memory leak.", UserWarning

```
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 2.0min
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 9.0min
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 17.1min finished
```

```
[264]: RandomizedSearchCV(cv=3, error_score=nan,
                        estimator=RandomForestClassifier(bootstrap=True,
                                                            ccp_alpha=0.0,
                                                            class_weight=None,
                                                            criterion='gini',
                                                            max_depth=None,
                                                            max_features='auto',
                                                            max_leaf_nodes=None,
                                                            max_samples=None,
                                                            min_impurity_decrease=0.0,
                                                            min_impurity_split=None,
                                                            min_samples_leaf=1,
                                                            min_samples_split=2,
                                                            min_weight_fraction_leaf=0.0,
                                                            n_estimators=100,
                                                            n_jobs=...
                        param_distributions={'bootstrap': [True, False],
                                             'max_depth': [10, 20, 30, 40, 50, 60,
                                                            70, 80, 90, 100, 110,
                                                            None],
                                             'max_features': ['auto', 'sqrt'],
                                             'min_samples_leaf': [1, 2, 4],
                                             'min_samples_split': [2, 5, 10],
                                             'n_estimators': [200, 400, 600, 800,
                                                            1000, 1200, 1400, 1600,
                                                            1800, 2000]},
                        pre_dispatch='2*n_jobs', random_state=42, refit=True,
                        return_train_score=False, scoring=None, verbose=2)
```

```
[265]: random_para.best_params_
```

```
[265]: {'n_estimators': 400,
        'min_samples_split': 10,
        'min_samples_leaf': 4,
        'max_features': 'auto',
        'max_depth': 70,
        'bootstrap': True}
```

```
[278]: random_2 = RandomForestClassifier(n_estimators=1400,min_samples_split=
        ↳10,min_samples_leaf= 2,max_features = 'sqrt',max_depth=80,bootstrap= True)
random_2.fit(train_x,train_y)
random_2_y = random_2.predict(test_x)
print(accuracy_score(random_2_y,test_y))
```

0.8621783574198096

Feature Selection

```
[267]: feature =
        ↳SelectFromModel(RandomForestClassifier(n_estimators=1400,min_samples_split=
        ↳10,min_samples_leaf= 2,max_features = 'sqrt',max_depth=80,bootstrap= True))
feature.fit(train_x,train_y)
feature_support = feature.get_support()
feature_selected = train_x.loc[:,feature_support].columns.tolist()
print(str(len(feature_selected)), 'selected features')
```

3 selected features

```
[268]: feature_selected
```

```
[268]: ['Age', 'Balance', 'NumOfProducts']
```

```
[269]: train_x_feature = train_x[['Age', 'Balance', 'NumOfProducts']]
train_x_feature.head()
```

```
[269]:
```

	Age	Balance	NumOfProducts
2371	0.019520	-0.136437	1
6521	0.413914	0.155280	2
9487	0.272773	-1.211056	1
2253	0.697698	-1.211056	1
2120	0.957958	0.655888	1

```
[270]: test_x_feature = test_x[['Age', 'Balance', 'NumOfProducts']]
test_x_feature.head()
```

```
[270]:
```

	Age	Balance	NumOfProducts
1624	0.056557	-1.211056	2
2169	0.766266	-0.022564	2
4285	0.655155	0.663712	1

71	0.155155	-1.211056	1
4049	0.318318	-1.211056	2

```
[271]: random_2.fit(train_x_feature,train_y)
random_2_feature_y = random_2.predict(test_x_feature)
print(accuracy_score(random_2_feature_y,test_y))
```

0.8375044060627423

8 Naive Bayes

```
[272]: bayes = GaussianNB()
bayes.fit(train_x,train_y)
bayes_y = bayes.predict(test_x)
print(accuracy_score(bayes_y,test_y))
```

0.8463165315474093

Feature Selection

```
[273]: train_x_feature = train_x[["Age", "Balance"]] #based on correlation values
train_x_feature.head()
```

```
[273]:
```

	Age	Balance
2371	0.019520	-0.136437
6521	0.413914	0.155280
9487	0.272773	-1.211056
2253	0.697698	-1.211056
2120	0.957958	0.655888

```
[274]: test_x_feature = test_x[["Age", "Balance"]] #based on correlation values
test_x_feature.head()
```

```
[274]:
```

	Age	Balance
1624	0.056557	-1.211056
2169	0.766266	-0.022564
4285	0.655155	0.663712
71	0.155155	-1.211056
4049	0.318318	-1.211056

```
[275]: bayes.fit(train_x_feature,train_y)
bayes_feature_y =bayes.predict(test_x_feature)
print(accuracy_score(bayes_feature_y, test_y))
```

0.8121254846669017

9 Conclusion

The highest accuracy we achieved is by hypertuning RandomForestClassifier and using all the features.

```
[276]: print(str((accuracy_score(random_2_y,test_y)) * 100) + "%")
```

```
86.25308424391963%
```

When we get time this notebook will be updated

If you like my work, please upvote.